


Introduction to Stochastic Processes

MATH 310-2

© Jason Sieffen, 2016
Creative Commons By-Attribution Share-Alike 

Introduction

In Math 310-1, you studied random variables and sequences of random variables that were *independent*. Because they were independent, they satisfied properties like the central limit theorem and the law of large numbers.

But, many worldly phenomena are not independent. In this class we will study a basic type of non-independent process—Markov chains.

Markov Chain

DEFINITION

A sequence of random variables (X_n) is a **Markov chain** (or Markov for short) if for n and any sequence of events (A_i) having positive probability,

$$\mathbb{P}(X_n \in A_n \mid X_{n-1} \in A_{n-1}) = \mathbb{P}(X_n \in A_n \mid X_{n-1} \in A_{n-1}, X_{n-2} \in A_{n-2}, \dots).$$

In other words, (X_n) is *Markov* if knowing the previous value in the sequence gives you the same amount of information as knowing all previous values. Since we will be looking at conditional probabilities a lot, to keep things less cluttered, we will sometimes avoid writing the events. That is,

$$\mathbb{P}(X_n \in A_n \mid X_{n-1}, X_{n-2}, \dots) \equiv \mathbb{P}(X_n \in A_n \mid X_{n-1} \in A_{n-1}, X_{n-2} \in A_{n-2}, \dots)$$

where (A_i) is a sequence of positive probability events. Further, $\mathbb{P}(X_n)$, will mean the distribution of X_n .

Let's look at some sequences and decide whether or not they are Markov.

1. You're walking around Chicago. At each street intersection, you roll a die to decide which street to go down next. (X_n) is the sequence of streets you walk down.

Markov! The next street you decide to go down depends only on your current position, not on how you got there.

2. You're trying to hack a computer network. Every minute you have a 50% chance to crack another password. If you've cracked $k/2$ passwords in the last k minutes, your heart-rate gets high, otherwise it is normal. (X_i) is your heart-rate in the i th minute.

Markov, but not *stationary*. For minute n , if you are told that X_{n-1} = caught, you know with probability 1 that X_n will be caught. Alternatively, if you're told X_{n-1} = not caught, you know that X_{n-2}, \dots = not caught, and so you get know extra information knowing all previous states.

3. You're trying to hack a computer network. Every minute you have a 50% chance to crack another password. If you crack k passwords, you will be discovered. (X_i) is how many passwords you've cracked by the i th minute.

Markov! If you've cracked n passwords in the i th minute, you still have a 50/50 chance of cracking your $(n+1)$ th password in the $(i+1)$ th minute, regardless of whether you cracked n or $n-1$ in the $(i-1)$ th minute.

4. (X_i) is an independent and identically distributed (iid) sequence of Bernoulli $(1/3, 2/3)$ random variables.

Markov! We trivially have $\mathbb{P}(X_i \mid X_{i-1}) = \mathbb{P}(X_i) = \mathbb{P}(X_i \mid X_{i-1}, X_{i-2}, \dots)$ by independence.

5. You repeatedly take a test to become part of your favorite fan club. (X_i) is your average score on the first i tests.

Markov, but not *stationary*. This one is subtle. Suppose that your score on the i th test is given by the independent sequence S_i . Then

$$X_i = \frac{1}{n} \sum_{n=1}^i S_i.$$

We then have $\mathbb{P}(X_i = a \mid X_{i-1} = b) = \mathbb{P}\left(\frac{S_i + (i-1)b}{i} = a\right)$. Since S_i is an independent sequence, knowing X_{i-2} doesn't change this probability.

Day 2

Notation

Let (X_i) have the Markov property. We call the codomain $(X_i) = X^{(i)}$ the *state space*, and we will assume $X^{(i)} = X^{(j)}$.

$\mathbb{P}(X_i = b \mid X_{i-1} = a)$ is called the *transition probability* from state a to state b at time i . If

$$\mathbb{P}(X_i = b \mid X_{i-1} = a) = \mathbb{P}(X_j = b \mid X_{j-1} = a) \quad \text{for all } i \text{ and } j,$$

we call (X_i) *stationary*. We will only consider stationary processes.

For a stationary process (X_i) ,

$$P(a, b) \equiv \mathbb{P}(X_i = b \mid X_{i-1} = a) \quad \text{and} \quad P^n(a, b) \equiv \mathbb{P}(X_i = b \mid X_{i-n} = a)$$


are called the *one-step transition function* and the *multi-step transition function*.

A Simple Example

Your very conservative aunt decided to go to a Chicago comedy club. Each joke leaves her happy (H) or angry (A). If she is happy, the next joke has a 50/50 chance of leaving her happy or angry. If she is angry, the next joke has a 1/4th chance of making her happy and a 3/4th chance of making her angry.

Day 3

We can model this situation with a directed graph:



DEF

The **transition matrix** associated with a stationary Markov chain (X_i) having n states $\{1, 2, \dots, n\}$ is the matrix $T = [t_{ij}]$ where $t_{ij} = P(i, j)$.

For this example, letting happy be the first state and angry be the second state, we have

$$T = \begin{matrix} & \begin{matrix} H & A \end{matrix} \\ \begin{matrix} H \\ A \end{matrix} & \begin{bmatrix} \mathbb{P}(H \rightarrow H) & \mathbb{P}(H \rightarrow A) \\ \mathbb{P}(A \rightarrow H) & \mathbb{P}(A \rightarrow A) \end{bmatrix} \end{matrix} = \begin{bmatrix} 1/2 & 1/2 \\ 1/4 & 3/4 \end{bmatrix}.$$

Let's compute $P^2(A, A) = \mathbb{P}(X_i = A \mid X_{i-2} = A)$. There are only two ways for this to happen. Either the states went

$$A \rightarrow A \rightarrow A \quad \text{or} \quad A \rightarrow H \rightarrow A.$$

$\mathbb{P}(A \rightarrow A \rightarrow A) = \frac{3}{4} \cdot \frac{3}{4} = \frac{9}{16}$ and $\mathbb{P}(A \rightarrow H \rightarrow A) = \frac{1}{4} \cdot \frac{1}{2} = \frac{2}{16}$, and so

$$P^2(A, A) = \mathbb{P}(A \rightarrow A \rightarrow A) + \mathbb{P}(A \rightarrow H \rightarrow A) = \frac{11}{16}.$$

The P^2 transition matrix now looks like $\begin{bmatrix} ? & ? \\ ? & \frac{11}{16} \end{bmatrix}$. Computing every entry we get $\begin{bmatrix} \frac{6}{16} & \frac{10}{16} \\ \frac{9}{16} & \frac{11}{16} \end{bmatrix}$.

Let's look at what happened from a matrix perspective: To compute $P^2(A, A)$, we need to compute the quantity $Q = \mathbb{P}(A \rightarrow A \rightarrow A) + \mathbb{P}(A \rightarrow H \rightarrow A)$, and by the Markov property,

$$\mathbb{P}(A \rightarrow A \rightarrow A) = \mathbb{P}(A \rightarrow A)\mathbb{P}(A \rightarrow A) \quad \text{and} \quad \mathbb{P}(A \rightarrow H \rightarrow A) = \mathbb{P}(A \rightarrow H)\mathbb{P}(H \rightarrow A).$$

If we squint, we see

$$Q = \begin{bmatrix} \mathbb{P}(H \rightarrow A) \\ \mathbb{P}(A \rightarrow A) \end{bmatrix} \cdot \begin{bmatrix} \mathbb{P}(A \rightarrow H) \\ \mathbb{P}(A \rightarrow A) \end{bmatrix} = \begin{bmatrix} \mathbb{P}(H \rightarrow A) & \mathbb{P}(A \rightarrow A) \end{bmatrix} \begin{bmatrix} \mathbb{P}(A \rightarrow H) \\ \mathbb{P}(A \rightarrow A) \end{bmatrix} = \begin{bmatrix} 0 & 1 \end{bmatrix} T \begin{pmatrix} T \begin{bmatrix} 0 \\ 1 \end{bmatrix} \end{pmatrix} = \begin{bmatrix} 0 & 1 \end{bmatrix} T^2 \begin{bmatrix} 0 \\ 1 \end{bmatrix}.$$

In other words, the $(2, 2)$ entry of the transition matrix for $P^2(x, y)$ is the $(2, 2)$ entry of T^2 . In fact, the transition function $P^2(x, y)$ is given by the matrix T^2 !

Using induction, we can show that the transition function $P^n(x, y)$ is given by the matrix T^n for any n , but first let's prove a lemma we've already used once.

LEMMA

If (X_i) is a Markov chain and (s_i) is a sequence of states of (X_i) , then

$$\mathbb{P}(s_1 \rightarrow s_2 \rightarrow \dots \rightarrow s_n) = \prod_{i=1}^{n-1} \mathbb{P}(s_i \rightarrow s_{i+1}).$$

We used this lemma already when we said $\mathbb{P}(A \rightarrow H \rightarrow A) = \mathbb{P}(A \rightarrow H)\mathbb{P}(H \rightarrow A)$, and it seems quite intuitive given the definition of Markov.

Proof. We will proceed by induction. Trivially, $\mathbb{P}(s_1 \rightarrow s_2) = \mathbb{P}(s_1 \rightarrow s_2)$. Suppose $\mathbb{P}(s_1 \rightsquigarrow s_n) = \prod \mathbb{P}(s_i \rightarrow s_{i+1})$, and consider $\mathbb{P}(s_1 \rightsquigarrow s_n \rightarrow s_{n+1})$. By definition

$$\begin{aligned} \mathbb{P}(s_n \rightarrow s_{n+1}) &= \mathbb{P}(X_{n+1} = s_{n+1} \mid X_n = s_n) \\ &= \mathbb{P}(X_{n+1} = s_{n+1} \mid X_n = s_n, X_{n-1} = s_{n-1}, \dots) \end{aligned}$$

with the last equality following from the Markov property. Now by the definition of conditional probability we have

$$\mathbb{P}(X_{n+1} = s_{n+1} \mid X_n = s_n, X_{n-1} = s_{n-1}, \dots) = \frac{\mathbb{P}(X_{n+1} = s_{n+1}, X_n = s_n, X_{n-1} = s_{n-1}, \dots)}{\mathbb{P}(X_n = s_n, X_{n-1} = s_{n-1}, \dots)} = \frac{\mathbb{P}(s_1 \rightsquigarrow s_n \rightarrow s_{n+1})}{\mathbb{P}(s_1 \rightsquigarrow s_n)}.$$

Multiplying both sides by $\mathbb{P}(s_1 \rightsquigarrow s_n)$ gives us

$$\mathbb{P}(s_1 \rightsquigarrow s_n \rightarrow s_{n+1}) = \mathbb{P}(s_1 \rightsquigarrow s_n)\mathbb{P}(s_n \rightarrow s_{n+1}).$$

Since we assume that $\mathbb{P}(s_1 \rightsquigarrow s_n)$ could be written as a product, we've completed the proof. \square

Day 4

We're almost ready to tackle our first theorem, but let's first prove two more easy facts.

LEMMA

If (X_i) is a Markov chain with finite state space S , then

$$\mathbb{P}(X_n = b) = \sum_{s \in S} \mathbb{P}(X_n = b \mid X_{n-1} = s)\mathbb{P}(X_{n-1} = s).$$

Proof. Since $\mathbb{P}(X_i \in S) = 1$, we have

$$\mathbb{P}(X_n = b) = \mathbb{P}(X_n = b \text{ and } X_{n-1} \in S) = \mathbb{P}\left(\bigcup_{s \in S} X_n = b \text{ and } X_{n-1} = s\right).$$

Since the events $(X_n = b \text{ and } X_{n-1} = s)$ and $(X_n = b \text{ and } X_{n-1} = s')$ are disjoint if $s \neq s'$, we have

$$\mathbb{P}\left(\bigcup_{s \in S} X_n = b \text{ and } X_{n-1} = s\right) = \sum_{s \in S} \mathbb{P}(X_n = b \text{ and } X_{n-1} = s).$$

Noticing the summand is just $\mathbb{P}(X_n = b \mid X_{n-1} = s)\mathbb{P}(X_{n-1} = s)$ completes the proof. \square

LEM

If (X_i) is a Markov chain with finite state space S , then $\mathbb{P}(X_n = b \mid X_1 = a)$ is the sum of $\mathbb{P}(a \rightarrow s_2 \rightarrow s_3 \rightarrow \dots \rightarrow s_{n-1} \rightarrow b)$ for every possible combination of $s_2, \dots, s_{n-1} \in S$.

Proof. Like the previous proof, this follows from manipulations of the basic definitions.

First note, if s_2, \dots, s_{n-1} and s'_2, \dots, s'_{n-1} are not the exact same sequence, then the events $a \rightarrow s_2 \rightarrow \dots \rightarrow s_{n-1} \rightarrow b$ and $a \rightarrow s'_2 \rightarrow \dots \rightarrow s'_{n-1} \rightarrow b$ are disjoint. Thus

$$\mathbb{P}\left(\bigcup_{s_2, \dots, s_{n-1} \in S} a \rightarrow s_2 \rightarrow \dots \rightarrow s_{n-1} \rightarrow b\right) = \sum_{s_2, \dots, s_{n-1} \in S} \mathbb{P}(a \rightarrow s_2 \rightarrow \dots \rightarrow s_{n-1} \rightarrow b).$$

Now, by the definition of conditional probability,

$$\mathbb{P}(a \rightarrow s_2 \rightarrow \dots \rightarrow s_{n-1} \rightarrow b) = \frac{\mathbb{P}(X_1 = a, X_2 = s_2, \dots, X_{n-1} = s_{n-1}, X_n = b)}{\mathbb{P}(X_1 = a)},$$

and so

$$\mathbb{P}\left(\bigcup_{s_2, \dots, s_{n-1} \in S} a \rightarrow s_2 \rightarrow \dots \rightarrow s_{n-1} \rightarrow b\right) = \frac{\mathbb{P}(X_1 = a, X_2 \in S, \dots, X_{n-1} \in S, X_n = b)}{\mathbb{P}(X_1 = a)} = \frac{\mathbb{P}(X_1 = a, X_n = b)}{\mathbb{P}(X_1 = a)},$$

with the last equality following because $X_i \in S$ always. The right side of this equation is just $\mathbb{P}(X_n = b \mid X_1 = a)$, and so combining all of our equations proves the result. \square

Now we're ready to tackle our theorem linking matrix multiplication and multi-step transition functions.

THEOREM

For a Markov chain (X_i) having states $\{1, \dots, d\}$ and transition function $P(x, y)$, let T be the corresponding transition matrix. Then, the transition matrix corresponding to the multi-step transition function $P^n(x, y)$ is T^n .

Proof. We will proceed by induction. By definition, T gives the transition probabilities for $P(x, y)$. Suppose that T^{n-1} gives the transition probabilities for $P^{n-1}(x, y)$. Since (X_i) is a Markov processes,

$$\mathbb{P}(s_1 \rightsquigarrow s_{n-1} \rightarrow s_n) = \mathbb{P}(s_1 \rightsquigarrow s_{n-1})\mathbb{P}(s_{n-1} \rightarrow s_n),$$

where s_1, \dots, s_n are states. Further, from our previous lemmas,

$$\begin{aligned} \mathbb{P}(X_n = s_n \mid X_1 = s_1) &= \sum_{s \in S} \mathbb{P}(X_n = s_n \mid X_{n-1} = s_{n-1})\mathbb{P}(X_{n-1} = s_{n-1} \mid X_1 = s_1) \\ &= \sum_{s \in S} P(s, s_n)P^{n-1}(s_1, s) = \sum_{s \in S} P^{n-1}(s_1, s)P(s, s_n). \end{aligned}$$

However, the right hand side is the same as the dot product of the s_n column of T with the s_1 row of T^{n-1} , which is the definition of the (s_1, s_n) entry of $TT^{n-1} = T^n$. \square

Now we know how to analyze transition probabilities for any finite-state Markov chain using matrices!

Let's see how our probability questions become linear algebra questions. Suppose we'd like to know the distribution after one joke supposing that your aunt started happy. That is, we'd like to know $f(x) = P(H, x)$. This is just the first row of T , and so the distribution is given by $\begin{bmatrix} 1 & 0 \end{bmatrix} T$. The distribution after 4 steps supposing your aunt started off angry is $\begin{bmatrix} 0 & 1 \end{bmatrix} T^7$. And $P^7(A, H)$ is $\begin{bmatrix} 0 & 1 \end{bmatrix} T^7 \begin{bmatrix} 1 \\ 0 \end{bmatrix}$.

But we can do more. Suppose we wanted to know the distribution after one joke supposing that there was a 50/50 chance your aunt came into the club happy or angry. That's just $\begin{bmatrix} 1/2 & 1/2 \end{bmatrix} T$. In general,

if \vec{d} is a row vector specifying an initial distribution, then $\vec{d}T^n$ gives the resulting distribution after n steps.

Day 5

Consider now the following questions: After many many jokes, what is the probability that your aunt leaves the club happy? What's the probability she leaves angry? Does the state she comes into the club with affect these probabilities? How does the number of jokes told affect these probabilities?

We can answer all of these questions with linear algebra! Since we're going to be computing large powers of T , let's diagonalize it.

$$T = \begin{bmatrix} 1/2 & 1/2 \\ 1/4 & 3/4 \end{bmatrix} = \begin{bmatrix} -2 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 1/4 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} -1/3 & 1/3 \\ 1/3 & 2/3 \end{bmatrix} = SDS^{-1}.$$

Since $(1/4)^n \rightarrow 0$ as $n \rightarrow \infty$, we can easily compute

$$T^\infty \equiv \lim_{n \rightarrow \infty} T^n = \begin{bmatrix} -2 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} -1/3 & 1/3 \\ 1/3 & 2/3 \end{bmatrix} = \begin{bmatrix} 1/3 & 2/3 \\ 1/3 & 2/3 \end{bmatrix}.$$

From this we can see $\begin{bmatrix} 1 & 0 \end{bmatrix} T^\infty = \begin{bmatrix} 0 & 1 \end{bmatrix} T^\infty$ and so her chances of being happy or angry upon leaving the club after many jokes doesn't depend on her initial state. In either case, she leaves happy with a $1/3$ chance and angry with a $2/3$ chance. We can also conclude that the number of jokes told doesn't really matter as long as the number is large. Indeed

$$T^n - T^m = S(D^n - D^m)S^{-1} = S \begin{bmatrix} \frac{1}{4^n} - \frac{1}{4^m} & 0 \\ 0 & 0 \end{bmatrix} S^{-1}.$$

As long as n and m are large, $\frac{1}{4^n} - \frac{1}{4^m} \approx 0$, and so the difference is approximately the zero matrix.

Day 6

Hitting Times

Consider a Markov chain for a gambler. The state space would be $\$d \geq 0$. If the gambler has any money, he can bet it and transition to a state with more money or a state with less money. But, if the gambler ever has $\$0$, he cannot bet any further.

The *Gambler's Ruin* problem asks, how long does it take to end up in state $\$0$?

Hitting Time

Let (X_i) be a Markov chain. A **hitting time** for an event B is the random variable

$$T_B = \min\{i \geq 0 : X_i \in B\}.$$

We will not compute hitting times right now, but hitting times give us another way to think about Markov chains. For example,

$$P^n(x, y) = \sum_{m \leq n} P(T_{\{y\}} = m \mid X_1 = x) P^{n-m}(y, y).$$

(Why is this formula intuitive?)

Different Types of States

Consider the gambler's ruin problem. If $a, b > 0$, there's a non-zero chance of starting at state a and ending at state b . (Is the correct way to mathematical quantify this statement $P(a, b) > 0$?) However, $P^n(0, b) = 0$ unless $b = 0$. That is, you cannot transition out of state $\$0$.

Absorbing

A state a of a Markov chain is called **absorbing** if $P(a, b) = 0$ whenever $b \neq a$.

In plain language, a state is absorbing if once you enter it, you can never leave.

For a Markov chain, define

$$\rho_{xy} = \mathbb{P}(T_{\{y\}} < \infty \mid X_1 = x).$$

That is, ρ_{xy} is the probability you end up at state y at some time in the future having started at state x .

Recurrent & Transient

For a Markov chain (X_i) , a state a is called **recurrent** if $\rho_{aa} = 1$, otherwise it is called **transient**.

What are the transient and recurrent states for the gambler's ruin problem? Well, the state \$0 is absorbing, and if the odds are in the house's favor, $\rho_{x0} = 1$. That is, you will always eventually end up at \$0. Thus, the state \$0 is recurrent. However, if you start at state \$d > 0, there is a positive probability you end up at state \$0, and so $\rho_{dd} < 1$. This means that every state \$d > 0 is transient.

Now we get to our first big theorem.

Let (X_i) be a Markov chain, let $N(x)$ be the number of times state x is visited, and let $G(y, x) = \mathbb{E}(N(x) \mid X_1 = y)$. Then,

(i) if x is recurrent, $\mathbb{P}(N(x) = \infty \mid X_1 = x) = 1$ and $G(x, x) = \infty$, and

(ii) if x is transient, $\mathbb{P}(N(x) < \infty \mid X_1 = y) = 1$ and

$$G(y, x) = \frac{\rho_{yx}}{1 - \rho_{xx}}.$$

In plain words, the theorem says that recurrent states are visited infinitely often and transient states are visited finitely often. Also, we can precisely compute the expectation of the number of times a transient state is visited in terms of ρ .

Proof. (i). Suppose x is a recurrent state. Let $T_{\{x\}}^{(n)} = \min\{i > n : X_i = x\}$ be the first return time to $\{x\}$ after time n . Since X_i is a stationary process,

$$\mathbb{P}(T_{\{x\}} < \infty \mid X_1 = x) = \mathbb{P}(T_{\{x\}}^{(n)} < \infty \mid X_n = x).$$

Since x is recurrent, $\mathbb{P}(T_{\{x\}} < \infty \mid X_1 = x) = 1$, and so $\mathbb{P}(T_{\{x\}}^{(n)} < \infty \mid X_n = x) = 1$. Thus if we are ever in state x we must again be in state x at some point in the future, and so $N(x) = \infty$. It immediately follows that $G(x, x) = \mathbb{E}(\infty) = \infty$.

(ii). Suppose x is transient. Let p_n be the probability that, starting at y , we hit x at least n times. Computing with the Markov property, we see

$$p_n = \rho_{yx} \rho_{xx}^{n-1} = \mathbb{P}(N(x) \geq n \mid X_1 = y).$$

We'd like to compute $\mathbb{P}(N(x) < \infty \mid X_1 = y)$. Applying an easy lower bound, we see

$$\begin{aligned} \mathbb{P}(N(x) < \infty \mid X_1 = y) &= \sum_{i < \infty} \mathbb{P}(N(x) = i \mid X_1 = y) \\ &\geq \sum_{i \leq n} \mathbb{P}(N(x) = i \mid X_1 = y) \\ &= 1 - \mathbb{P}(N(x) \geq n \mid X_1 = y) = 1 - \rho_{yx} \rho_{xx}^{n-1}, \end{aligned}$$

and so $\mathbb{P}(N(x) < \infty \mid X_1 = y) = 1$.

Now, let $d(n) = \mathbb{P}(N(x) = n \mid X_1 = y)$ be the distribution of $N(x)$. We can explicitly write d as

$$\begin{aligned} d(n) &= \mathbb{P}(N(x) \geq n \mid X_1 = y) - \mathbb{P}(N(x) \geq n+1 \mid X_1 = y) \\ &= \rho_{yx} \rho_{xx}^{n-1} - \rho_{yx} \rho_{xx}^n = \rho_{yx} \rho_{xx}^{n-1} (1 - \rho_{xx}), \end{aligned}$$

and so

$$G(y, x) = \mathbb{E}(N(x) \mid X_1 = y) = \sum n d(n) = \sum n \rho_{yx} \rho_{xx}^{n-1} (1 - \rho_{xx}) = \frac{\rho_{yx}}{1 - \rho_{xx}}.$$

□